Comparative Review of DNA and RNA Sequence Analysis Tools in Modern Bioinformatics

Amir Mahdi Mahboubi
Bachelor Student of Computer Engineering

Email: amirmahdi.mahboubi@iau.ir
Alternative Email: amir.mahdi.mahboubi2005@gmail.com
ORCID: 0009-0007-8924-4421
LinkedIn: www.linkedin.com/in/amir-mahdi-mahboubi-521542306

Abstract

Sequence analysis lies at the computational core of contemporary molecular biology, translating raw nucleotide reads into biological knowledge that informs genomics, transcriptomics, and precision medicine. The proliferation of next-generation and third-generation sequencing platforms has been matched by an equally rapid expansion of software designed to align, assemble, quantify, and annotate DNA and RNA sequences, leaving researchers with a crowded and often confusing landscape of competing tools. This review provides a critical, comparative examination of the algorithmic foundations, computational performance, and practical trade-offs of the most widely used sequence analysis tools in current bioinformatics practice. DNA-oriented aligners, namely BLAST, BWA, Bowtie2, HISAT2, and Minimap2, are evaluated against criteria of indexing strategy, computational complexity, memory footprint, and suitability for short- versus long-read data. RNA-oriented tools, including STAR, Salmon, Kallisto, StringTie, and Cufflinks, are assessed with respect to alignment-based versus alignment-free quantification paradigms, isoform reconstruction accuracy, and runtime efficiency. The review further synthesizes evidence on genome assembly, variant calling, and functional annotation pipelines, and devotes particular attention to the accelerating integration of deep learning and transformer-based genomic language models, such as DNABERT, the Nucleotide Transformer, and Enformer, into conventional bioinformatics workflows. Persistent challenges, including the scaling of computational infrastructure to petabyte-level genomic datasets, reproducibility across heterogeneous pipelines, and ethical concerns surrounding genomic privacy, are discussed alongside emerging directions in cloud-native bioinformatics, multi-omics integration, and real-time genomic analytics. The review concludes with practical, evidence-based recommendations to guide tool selection according to study design, data modality, and available computational resources, offering a reference point for students and early-career researchers entering computational genomics.

Keywords: bioinformatics; DNA sequencing; RNA-seq; sequence alignment; genome assembly; variant calling; transcriptomics; deep learning genomics

1. Introduction

Bioinformatics emerged as a discipline at the intersection of biology, computer science, and statistics in response to a fundamental problem: biological sequence data accumulates far faster than unaided human reasoning can interpret it [1]. What began in the 1970s as a modest exercise in protein sequence comparison has matured into a field that underwrites nearly every branch of modern life science, from evolutionary biology to oncology. The discipline's growth has been propelled less by theoretical biology than by computational necessity, since the volume, dimensionality, and noise characteristics of genomic data demand algorithmic solutions that classical wet-lab methods cannot provide.

Sequence analysis, the systematic computational interpretation of DNA and RNA sequences, sits at the center of this enterprise. It encompasses alignment of reads to reference genomes, de novo assembly of genomes lacking a reference, quantification of transcript abundance, identification of genetic variants, and functional annotation of genes and regulatory elements. Each of these tasks has historically been addressed by a distinct generation of software tools, and the choice among them materially affects downstream biological conclusions [2], [3]. A poorly chosen aligner can systematically miscall variants; an inappropriate transcript quantification method can distort differential expression results; an assembler ill-suited to a genome's repeat structure can fragment an assembly beyond usefulness. Tool selection is therefore not a peripheral technical decision but a methodological one with direct bearing on scientific validity.

The evolution from traditional, low-throughput biological analysis to computational genomics has occurred in two broad waves. The first wave, spanning Sanger sequencing and early capillary electrophoresis platforms, supported sequence comparison at the scale of individual genes or small genomes and was adequately served by exhaustive dynamic-programming algorithms such as

Needleman-Wunsch and Smith-Waterman [4]. The second wave, inaugurated by next-generation sequencing (NGS) platforms in the mid-2000s, increased throughput by several orders of magnitude while simultaneously shortening individual read lengths, a combination that rendered exhaustive alignment computationally infeasible and catalyzed the development of index-based heuristic aligners such as BWA and Bowtie [5], [6]. A third wave, driven by single-molecule long-read platforms from Pacific Biosciences and Oxford Nanopore Technologies, has more recently reintroduced long contiguous reads but at substantially higher per-base error rates, again reshaping the algorithmic requirements of alignment and assembly software [7], [8].

The significance of robust sequence analysis tools extends well beyond academic genomics. Clinical diagnostics increasingly rely on rapid, accurate variant calling from whole-exome and whole-genome sequencing data to identify pathogenic mutations [9]. Agricultural genomics depends on efficient assembly and annotation pipelines to accelerate crop and livestock breeding programs. Public health surveillance of viral and bacterial pathogens, a capability whose importance was made starkly visible during recent global outbreaks, depends on fast, reliable alignment and variant-calling software capable of processing samples within clinically meaningful timeframes. Against this backdrop, the present review undertakes a systematic, comparative assessment of the sequence analysis tools that dominate contemporary practice, with attention to their underlying algorithms, computational trade-offs, and appropriate use cases, and situates this assessment within the broader, rapidly evolving context of machine learning-augmented genomics.

The remainder of this paper is organized as follows. Section 2 reviews the fundamentals of DNA and RNA sequence analysis and the sequencing technologies that generate the underlying data. Section 3 categorizes the principal classes of sequence analysis software. Sections 4 and 5 provide detailed comparative analyses of major DNA- and RNA-oriented tools, respectively, supported by comparative tables in Section 6. Section 7 examines the integration of machine learning and transformer-based models into sequence analysis. Sections 8 and 9 discuss persistent challenges and future directions, and Section 10 concludes with practical recommendations for tool selection.

2. Fundamentals of DNA and RNA Sequence Analysis

2.1 DNA Structure and Sequencing Concepts

Deoxyribonucleic acid (DNA) encodes genetic information as a linear sequence of four nucleotide bases, adenine, thymine, guanine, and cytosine, arranged along a double helix. Sequencing technologies determine this base order experimentally, but no current platform reads an entire chromosome contiguously; instead, sequencers generate fragmented reads that must be computationally reassembled or mapped to a known reference. This fragmentation is the central reason sequence analysis software exists at all, and the statistical and algorithmic properties of the fragmentation process, read length, per-base error rate, and coverage depth, determine which computational strategies are viable. Short-read platforms such as Illumina sequencing-by-synthesis produce reads of 100 to 300 base pairs with per-base error rates below 1 percent, favoring exact or near-exact index-based alignment algorithms [10]. Long-read platforms produce reads from several kilobases to over a megabase, with historically higher error rates, now substantially reduced by circular consensus and basecalling improvements, necessitating alignment algorithms tolerant of higher mismatch and indel rates [7].

2.2 RNA Structure and Transcriptomics

Ribonucleic acid (RNA), in particular messenger RNA (mRNA), constitutes the functional intermediate between the genome and the proteome and is the principal substrate of transcriptomic analysis. RNA sequencing (RNA-seq) measures the relative abundance of transcripts within a sample, enabling quantification of gene expression, detection of alternative splicing, and discovery of novel transcripts and fusion genes [11]. Unlike DNA, transcribed RNA in eukaryotes undergoes splicing, in which intronic sequences are excised and exons are joined, so that sequencing reads derived from mature mRNA do not align contiguously to the genomic reference but instead span exon-exon junctions. This splicing phenomenon is the principal reason RNA-seq alignment requires specialized splice-aware aligners distinct from those used for genomic DNA, and it is also the reason alignment-free quantification methods, which bypass genomic alignment entirely in favor of direct comparison against a transcriptome index, have gained substantial traction.

2.3 Next-Generation Sequencing

NGS technologies, epitomized by the Illumina platform, perform massively parallel sequencing of millions to billions of short DNA fragments simultaneously, reducing the cost per base by several orders of magnitude relative to Sanger sequencing [12]. This throughput gain came at the cost of read length and introduced a distinctive computational burden: aligning billions of short reads against reference genomes of gigabase scale within practical runtimes and memory budgets. The response of the bioinformatics community was the development of full-text index data structures, most notably the Burrows-Wheeler Transform (BWT) combined with the Ferragina-Manzini (FM) index, which permit exact and approximate substring matching in time roughly proportional to read length rather than reference size [5], [13]. These index structures underlie BWA, Bowtie2, and HISAT2, and remain the dominant approach for short-read alignment to the present day.

2.4 Third-Generation Sequencing Technologies

Third-generation, single-molecule sequencing platforms from Pacific Biosciences (PacBio) and Oxford Nanopore Technologies (ONT) sequence individual DNA or RNA molecules in real time without prior amplification, yielding reads that can span tens to hundreds of kilobases [8]. The introduction of PacBio HiFi sequencing, which generates circular consensus reads with accuracy exceeding 99.9 percent, has substantially narrowed the historical accuracy gap between long- and short-read technologies while retaining the structural advantages of long reads for resolving repetitive regions, structural variants, and full-length transcript isoforms [7]. Nanopore sequencing, while typically exhibiting somewhat higher raw error rates, offers the distinct advantages of ultra-long reads, real-time data streaming, and direct RNA sequencing without a reverse-transcription step, properties that have made it attractive for outbreak genomic surveillance and portable field sequencing [14].

2.5 Major Challenges in Sequence Analysis

Several challenges recur across the sequence analysis landscape regardless of platform. Repetitive genomic regions, which constitute roughly half of the human genome, produce ambiguous alignments that confound both short- and long-read aligners and complicate genome assembly [15]. Sequencing error profiles differ systematically across platforms, biasing variant calls unless the calling algorithm explicitly models the relevant error distribution. Computational scalability remains a persistent constraint, as whole-genome sequencing of large cohorts now routinely produces datasets in the multi-terabyte range, straining conventional computing infrastructure and motivating the migration of pipelines to cloud and high-performance computing environments [16]. Finally, reference bias, the tendency of alignment-based methods to underrepresent sequences divergent from the reference genome, has emerged as a recognized source of systematic error, motivating pangenome reference structures intended to better capture human genetic diversity.

3. Categories of Sequence Analysis Tools

Sequence analysis software can be organized into functional categories corresponding to distinct stages of the bioinformatics pipeline. Although individual tools sometimes span more than one category, this taxonomy provides a useful organizing framework for comparative evaluation.

3.1 Sequence Alignment Tools

Alignment tools map sequencing reads or full sequences to a reference, or to one another, identifying regions of similarity that reflect homology, conserved function, or simply identity to a known genome. Alignment algorithms divide broadly into exhaustive dynamic-programming methods, which guarantee optimal alignments but scale poorly with sequence length, and heuristic, index-based methods, which sacrifice the optimality guarantee for tractable runtime on large datasets [4], [5]. The dominant index structures in current practice are the suffix array and Burrows-Wheeler Transform, used by BWA and Bowtie2, and minimizer-based seeding, used by Minimap2, each offering different trade-offs between indexing time, memory footprint, and alignment sensitivity.

3.2 Genome Assembly Tools

Assembly tools reconstruct a genome sequence de novo from overlapping reads in the absence of, or in addition to, a reference sequence. Short-read assemblers predominantly use de Bruijn graph algorithms, which decompose reads into fixed-length k-mers and represent overlaps as graph edges, an approach exemplified by SPAdes [17]. Long-read assemblers instead favor overlap-layout-consensus or repeat-graph algorithms better suited to sparse, long, and error-prone reads, as implemented in Flye and hifiasm [18], [19]. Hybrid assembly strategies that combine short- and long-read data, such as Unicycler, attempt to capture the contiguity benefits of long reads alongside the base-level accuracy of short reads [20].

3.3 Transcriptome Analysis Tools

Transcriptome analysis tools quantify and characterize RNA transcripts, encompassing splice-aware genome alignment, alignment-free pseudoalignment and quantification, and transcript assembly. The methodological divide between alignment-based pipelines, such as STAR followed by StringTie, and alignment-free pipelines, such as Salmon or Kallisto, represents one of the more consequential design decisions in RNA-seq analysis, with direct implications for runtime, sensitivity to reference completeness, and downstream differential expression accuracy [21], [22].

3.4 Variant Calling Tools

Variant calling tools identify positions at which an individual's sequence differs from a reference, encompassing single-nucleotide polymorphisms, small insertions and deletions, and larger structural variants. Classical variant callers such as the Genome Analysis Toolkit (GATK) HaplotypeCaller rely on probabilistic, Bayesian models of sequencing error and haplotype likelihood [23], while more recent tools such as DeepVariant reformulate variant calling as an image classification problem solved by a convolutional neural network trained on pileup representations of aligned reads, achieving competitive or superior accuracy across sequencing platforms [24].

3.5 Functional Annotation Tools

Functional annotation tools assign biological meaning to genomic features, identifying gene models, regulatory elements, and the functional consequences of variants. Reference annotation projects such as GENCODE provide curated gene models against which novel transcripts and variants are interpreted [25], while annotation software such as SnpEff and ANNOVAR predict the functional impact of called variants on protein-coding sequence. Increasingly, functional annotation is being augmented by sequence-based deep learning models capable of predicting regulatory activity directly from DNA sequence without relying solely on existing annotation databases, a development discussed in Section 7.

4. Comparative Analysis of Major DNA Sequence Analysis Tools

This section examines five DNA-oriented alignment tools that, collectively, account for the substantial majority of alignment operations performed in contemporary genomics: BLAST, BWA, Bowtie2, HISAT2, and Minimap2. Each tool is assessed along consistent criteria, namely algorithmic basis, computational complexity, accuracy, runtime performance, memory requirements, and characteristic strengths, weaknesses, and applications.

4.1 BLAST

The Basic Local Alignment Search Tool (BLAST) remains, more than three decades after its introduction, the most widely recognized sequence comparison tool in biology [1]. BLAST employs a seed-and-extend heuristic: short exact-match seeds (default word size 11 for nucleotide searches) are identified between query and database sequences using a hash table, and seeds are subsequently extended using a simplified Smith-Waterman-like local alignment with a statistical significance threshold (E-value) governing which extensions are reported. This design trades the optimality guarantee of full dynamic programming for dramatically improved average-case runtime, with empirical complexity that scales sub-linearly with database size for typical queries owing to the indexing step [1], [26].

BLAST's principal strength lies not in raw throughput but in flexibility and interpretability: it accommodates highly divergent sequence comparisons across species, supports protein as well as nucleotide searches, and returns statistically calibrated significance scores that remain a standard reporting convention in comparative genomics. Its principal weakness is computational inefficiency at the scale of modern short-read sequencing; aligning hundreds of millions of 150-base-pair reads against a human genome with BLAST is computationally prohibitive relative to purpose-built short-read aligners, which is precisely why BLAST has been supplanted by BWA, Bowtie2, and their relatives for routine NGS read mapping. BLAST nonetheless retains an essential role in tasks for which those aligners are unsuited: homology searches against large, heterogeneous sequence databases, identification of novel or contaminant sequences, primer specificity checking, and comparative analysis of assembled contigs or genes across species. Memory requirements scale with database size but remain modest for typical gene-level or genome-level searches, and the BLAST+ reimplementation improved multithreading and I/O efficiency considerably over the legacy BLAST suite [26].

4.2 BWA

The Burrows-Wheeler Aligner (BWA) was among the first short-read aligners to exploit the FM-index, a compressed full-text index derived from the Burrows-Wheeler Transform, to achieve exact and mismatch-tolerant alignment in time roughly linear in read length and independent of reference genome size after indexing [5]. The original BWA-backtrack algorithm targeted reads up to roughly 100 base pairs, while the subsequently developed BWA-MEM algorithm extended applicability to longer Illumina reads and moderately long, moderately accurate reads through a seed-and-extend strategy based on maximal exact matches, followed by affine-gap Smith-Waterman extension [27].

BWA-MEM has become a de facto standard for clinical and population-scale short-read alignment, owing to a favorable combination of accuracy, speed, and, critically, broad compatibility with downstream variant-calling pipelines such as GATK, which were co-developed and extensively validated against BWA-MEM output [23]. Its computational complexity for seeding is close to linear in read length given the precomputed FM-index, though pathological cases involving highly repetitive regions can produce a combinatorial explosion of candidate seeds, mitigated in practice by seed-length and repeat-frequency heuristics. Memory requirements are dominated by the FM-index itself, which for the human genome occupies several gigabytes, a figure that has historically motivated the development of more memory-efficient successors for very large reference panels. BWA's principal limitation is that it is not splice-aware and is therefore unsuitable for direct RNA-seq alignment to a genomic reference, and its performance on long, error-prone third-generation reads is markedly inferior to aligners purpose-built for that data type, such as Minimap2.

4.3 Bowtie2

Bowtie2 likewise builds upon an FM-index but introduces a distinct seeding and extension strategy optimized for gapped alignment of reads from 50 base pairs up to several hundred base pairs in length [6]. Where the original Bowtie tool supported only ungapped alignment, Bowtie2 incorporates a dynamic-programming extension phase that permits insertions and deletions while retaining much of the speed advantage of pure index-based seeding. Bowtie2 additionally introduced multiseed alignment, in which several short seeds are extracted per read and aligned independently before consolidation, improving sensitivity to reads spanning structural variation relative to single-seed approaches.

In comparative benchmarks, Bowtie2 and BWA-MEM exhibit broadly similar accuracy and runtime for standard short-read genomic alignment, with the choice between them frequently determined by ecosystem compatibility rather than a clear performance differential; Bowtie2, for instance, remains the preferred aligner within several established ChIP-seq and ATAC-seq pipelines, while BWA-MEM dominates clinical variant-calling pipelines [28]. Bowtie2's memory footprint is comparable to BWA's, on the order of a few gigabytes for the human genome index, and its multithreaded implementation scales efficiently across modern multicore systems. Its principal limitation, shared with BWA, is the absence of native splice awareness, restricting direct applicability to RNA-seq without the specialized extensions implemented in tools such as TopHat, which has itself been largely superseded by HISAT2.

4.4 HISAT2

HISAT2 (Hierarchical Indexing for Spliced Alignment of Transcripts 2) extends the FM-index paradigm with a hierarchical graph-based index (FM-index combined with a hierarchical graph index, termed GFM) that simultaneously encodes the reference genome, known splice sites, and common genetic variants, enabling fast, splice-aware alignment of RNA-seq reads directly to the genome [29]. This hierarchical indexing strategy partitions the genome into many small, local FM-indexes anchored to a global index, allowing the aligner to efficiently identify candidate alignment locations across exon-exon junctions without the prohibitive memory cost of a single large suffix structure spanning the full search space of possible splice junctions.

HISAT2 is computationally efficient, commonly cited as among the fastest splice-aware aligners, processing tens of millions of reads per hour on modest hardware, with a memory footprint of approximately 4 to 6 gigabytes for the human genome, substantially lower than its predecessor TopHat2 and competitive with non-splice-aware aligners [29]. Its accuracy on known and novel splice junctions benefits from incorporation of annotated splice sites when available, though this also introduces a degree of reference-annotation dependency that can bias results toward previously characterized transcript structures. HISAT2 is typically paired with downstream transcript assembly tools such as StringTie, forming the HISAT2-StringTie pipeline that has become a standard, computationally efficient alternative to the STAR-based pipelines favored in many large-scale RNA-seq studies. Its principal weakness relative to STAR lies in somewhat reduced sensitivity for detecting novel, unannotated splice junctions in highly divergent or poorly annotated genomes.

4.5 Minimap2

Minimap2 was developed specifically to address the alignment demands of third-generation long-read sequencing data, employing a minimizer-based seeding strategy in which only a sparse, representative subset of k-mers per window is indexed, substantially reducing both index size and seeding time relative to exhaustive k-mer indexing [30]. Seeds are chained using a dynamic-programming-based co-linear chaining algorithm, and final base-level alignment is computed only where necessary using an efficient, SIMD-accelerated extension of the Smith-Waterman algorithm with concave gap penalties suited to the longer indels characteristic of long-read data.

Minimap2's principal strength is versatility combined with speed: a single tool, through preset parameter configurations, handles genomic DNA alignment, full-length mRNA or cDNA alignment with splice awareness, and assembly-to-assembly alignment, across both PacBio and Nanopore data, while remaining three to four times faster than mainstream short-read aligners at comparable accuracy on appropriate data and over thirty times faster than earlier dedicated long-read aligners [30]. Subsequent refinements improved alignment accuracy specifically for noisy long reads and short-read mapping scenarios alike [31]. Memory requirements are modest relative to long-read data volumes, generally a few gigabytes for human genome indexing, owing to the sparsity of the minimizer index. Minimap2's limitations are most apparent at the extremes: for very short reads below roughly 100 base pairs, dedicated short-read aligners retain a sensitivity advantage, and for applications demanding base-level alignment optimality rather than heuristic chaining, exhaustive methods remain preferable despite their computational cost. Nonetheless, Minimap2 has become the default aligner for nearly all long-read genomic and transcriptomic workflows, including those underlying recent telomere-to-telomere human genome assemblies.

5. Comparative Analysis of Major RNA Analysis Tools

RNA-seq analysis tools divide methodologically into alignment-based approaches, which map reads to a genome or transcriptome before quantification, and alignment-free approaches, which infer transcript abundance through lightweight pseudoalignment or quasi-mapping without computing base-level alignments. This section reviews STAR and Cufflinks as representatives of alignment-based pipelines, Salmon and Kallisto as representatives of alignment-free quantification, and StringTie as a transcript assembly and quantification tool bridging both paradigms.

5.1 STAR

STAR (Spliced Transcripts Alignment to a Reference) uses an uncompressed suffix array combined with a maximal mappable prefix search algorithm to identify splice junctions directly during the seeding phase, rather than relying on a precomputed annotation of splice sites, enabling sensitive detection of both known and novel junctions [32]. This approach allows STAR to align reads spanning long introns and complex splicing patterns with high sensitivity and to process reads at extremely high throughput, commonly cited as the fastest splice-aware aligner in standard benchmarks, owing to its highly parallelized, memory-resident suffix array search.

STAR's principal limitation is memory consumption: the human genome index typically requires approximately 27 to 32 gigabytes of RAM, an order of magnitude greater than HISAT2's hierarchical index, a constraint that can preclude its use on resource-limited infrastructure, although reduced-memory index-generation options partially mitigate this cost at some expense to runtime [32]. STAR is the alignment backbone of several major RNA-seq analysis ecosystems, including those used by large consortium projects, and is frequently paired downstream with quantification tools such as featureCounts or RSEM, or with Cufflinks-family assembly tools, and increasingly with STARsolo for single-cell RNA-seq alignment. Its accuracy in junction detection is generally regarded as best-in-class, particularly for discovery-oriented studies seeking novel isoforms, at the cost of substantially greater computational resource demand than alignment-free alternatives.

5.2 Salmon

Salmon performs alignment-free quantification through a quasi-mapping procedure that identifies the set of transcripts compatible with each read without computing a full base-level alignment, combined with a two-phase inference procedure that first estimates initial abundance via an expectation-maximization algorithm and then refines estimates using a variational Bayesian or online stochastic approach informed by a fragment-level generative model accounting for sequence-specific and GC-content biases [22]. By explicitly modeling these technical biases, Salmon aims to produce abundance estimates accurate enough to substantially reduce the need for separate downstream bias-correction steps.

Salmon's quantification speed is its defining advantage: typical RNA-seq samples are quantified in minutes on a single desktop-class machine, several orders of magnitude faster than alignment-based pipelines, with corresponding reductions in memory requirement, generally a few gigabytes for the human transcriptome index [22]. Salmon can additionally operate in a hybrid mode that accepts pre-computed genome alignments, partially bridging the alignment-based and alignment-free paradigms. Its principal limitation is a dependency on the completeness and accuracy of the reference transcriptome; because Salmon does not align to the genome, reads originating from unannotated transcripts, novel splice isoforms, or non-coding regions absent from the reference index cannot be quantified, a constraint absent from genome-alignment-based pipelines. Salmon's quantifications are widely used as direct input to differential expression tools such as DESeq2 and edgeR via the tximport interface, reflecting broad ecosystem integration.

5.3 Kallisto

Kallisto introduced the concept of pseudoalignment, in which reads are assigned to the equivalence class of transcripts with which they share compatible k-mers, determined via traversal of a de Bruijn graph constructed from the transcriptome, without computing explicit base-level alignments [21]. Transcript abundances are subsequently estimated from the resulting equivalence classes using an expectation-maximization algorithm closely related in spirit to that used by Salmon, though differing in implementation detail and bias-modeling sophistication.

Kallisto and Salmon are frequently regarded as functionally comparable in accuracy and are both dramatically faster than alignment-based quantification, with Kallisto's original publication reporting quantification of tens of millions of reads in a few minutes on standard hardware, a speed improvement of roughly two orders of magnitude over contemporary alignment-based pipelines [21]. Memory requirements are similarly modest, on the order of a few gigabytes for index construction and quantification of mammalian transcriptomes. Empirical comparisons in subsequent literature suggest Salmon's more detailed bias modeling can yield modestly improved accuracy under certain library preparation protocols, while Kallisto's simpler model can be marginally faster, though the practical difference between the two for most standard bulk RNA-seq applications is small [33]. Kallisto shares Salmon's principal limitation of transcriptome-reference dependency, and both tools require an additional step, such as kb-python or BUStools, to support droplet-based single-cell RNA-seq quantification, a use case for which Kallisto's pseudoalignment framework has been specifically extended.

5.4 StringTie

StringTie addresses transcript assembly and quantification given aligned RNA-seq reads, typically produced by HISAT2 or STAR, using a network flow algorithm to assemble the most parsimonious set of transcript isoforms consistent with observed splice junctions and coverage patterns, followed by maximum-flow-based abundance estimation for each assembled isoform [34]. Unlike Salmon and Kallisto, StringTie operates on genome alignments rather than a fixed transcriptome reference, allowing it to assemble and quantify novel, previously unannotated transcripts and isoforms directly from the data.

StringTie demonstrated improved accuracy in transcript reconstruction relative to its predecessor Cufflinks across multiple benchmark datasets, while also reducing runtime, a combination attributable to its more efficient flow-network formulation of the assembly problem [34]. The subsequent StringTie2 extension added explicit support for long-read RNA-seq data, accommodating the higher per-base error rates of PacBio and Nanopore platforms while preserving the core flow-based assembly logic, and improved handling of super-reads and mixed short- and long-read datasets [35]. StringTie's principal strength is its capacity for reference-guided transcript discovery, valuable in non-model organisms or disease contexts involving novel isoform usage, while its principal limitation relative to alignment-free quantification is substantially greater runtime and memory requirement, since it necessarily depends on a preceding full alignment step.

5.5 Cufflinks

Cufflinks, an earlier and historically influential transcript assembly and quantification tool, employs a different parsimony principle, constructing a minimum-path-cover over a graph of compatible reads and fragments to reconstruct the smallest set of transcript isoforms explaining the observed read alignments, with abundance subsequently estimated via a statistical model accounting for fragment length distribution and sequencing bias [36]. Cufflinks was historically paired with TopHat for alignment, forming the once-dominant Tuxedo pipeline, later extended to the Tuxedo2 pipeline incorporating HISAT2, StringTie, and Ballgown [34].

Cufflinks played a foundational role in establishing transcript-level, isoform-resolved RNA-seq analysis as a standard practice and remains historically significant for popularizing fragments-per-kilobase-per-million (FPKM) normalization [36]. However, subsequent benchmarking demonstrated that StringTie achieves higher transcript reconstruction accuracy and substantially faster runtime than Cufflinks on equivalent datasets, and the broader research community has largely transitioned toward StringTie or alignment-free quantification for new analyses [34].

Cufflinks remains relevant chiefly for reproducing or extending legacy analyses conducted under the original Tuxedo pipeline and for pedagogical purposes in illustrating the parsimony-based approach to isoform reconstruction, but is no longer recommended as a first choice for new RNA-seq projects given the demonstrated superiority of its successors on both accuracy and computational efficiency grounds.

6. Comparative Tables

The following tables summarize the comparative characteristics of the DNA- and RNA-oriented tools reviewed in Sections 4 and 5, synthesizing information on algorithmic basis, supported data types, accuracy, computational performance, and ecosystem maturity drawn from the original publications and subsequent independent benchmarking studies [1], [5], [6], [21], [22], [29], [30], [32], [34], [36], [37].

Table I. Comparative Overview of DNA Sequence Alignment Tools

Tool	Core Algorithm	Data Type	Typical Accuracy	Relative Speed	Memory (Human Genome)
BLAST	Seed-and-extend, local DP	Any length, cross-species	High (statistical E-value)	Low for NGS scale	Low-Moderate, DB-dependent
BWA-MEM	FM-index, MEM seeding	Short-moderate reads (Illumina)	Very high for short reads	High	~3-5 GB
Bowtie2	FM-index, multiseed	Short reads (50-500 bp)	High, comparable to BWA	High	~3-4 GB
HISAT2	Hierarchical FM/graph index	Splice-aware RNA + DNA	High, junction-sensitive	Very high	~4-6 GB
Minimap2	Minimizer seeding, co-linear chaining	Long reads, splice-aware, assembly	High for long/noisy reads	Very high	~3-6 GB

Table II. Comparative Overview of RNA Analysis Tools

Tool	Methodology	Quantification Strategy	Computational Efficiency	Key Use Case
STAR	Suffix array, MMP search	Counts via alignment + downstream tool	Fast, memory-intensive (~30 GB)	Novel junction discovery, single-cell
Salmon	Quasi-mapping	EM/VB with bias modeling	Very fast, low memory	Bulk RNA-seq quantification
Kallisto	Pseudoalignment (de Bruijn graph)	EM on equivalence classes	Very fast, low memory	Rapid quantification, scRNA-seq
StringTie	Network flow assembly	Flow-based isoform abundance	Moderate, alignment-dependent	Novel isoform discovery

Tool	Methodology	Quantification Strategy	Computational Efficiency	Key Use Case
Cufflinks	Minimum path cover	Statistical FPKM model	Slower, superseded by StringTie	Legacy Tuxedo pipeline

Table III. Scalability, Ease of Use, and Community Adoption

Tool	Scalability	Ease of Use	Community Adoption	Active Maintenance (as of 2026)
BLAST	Moderate (DB search scales sub-linearly)	High (web and CLI interfaces)	Extremely high, legacy standard	Yes (NCBI)
BWA-MEM	High (multithreaded)	Moderate (CLI, scripting required)	Extremely high, clinical standard	Yes
Bowtie2	High (multithreaded)	Moderate	High, ChIP-seq/ATAC-seq standard	Yes
HISAT2	High	Moderate	High, common RNA-seq pipeline	Yes
Minimap2	Very high	High (presets simplify use)	Very high, long-read standard	Yes
STAR	High but memory-bound	Moderate	Very high, consortium standard	Yes
Salmon	Very high	High	High, growing rapidly	Yes
Kallisto	Very high	High	High	Yes
StringTie	Moderate-High	Moderate	High	Yes
Cufflinks	Moderate	Moderate	Declining, legacy use	Limited

7. Machine Learning and AI in Sequence Analysis

7.1 Deep Learning Approaches

Deep learning has progressively displaced or augmented classical statistical methods across several stages of the sequence analysis pipeline. DeepVariant reframed variant calling as an image classification task, converting pileups of aligned reads around candidate variant sites into tensor representations and applying a convolutional neural network trained on benchmark truth sets to call genotypes, achieving accuracy that matched or exceeded the previously dominant GATK HaplotypeCaller across multiple sequencing technologies and notably generalizing across genome builds and species without retraining [24]. This generalization capacity, atypical of earlier rule-based variant callers, illustrates a broader trend in which learned representations transfer more readily across heterogeneous experimental conditions than hand-engineered statistical models.

Earlier convolutional architectures such as DeepSEA demonstrated that sequence-based deep models could predict chromatin accessibility and transcription factor binding directly from raw DNA sequence, without explicit feature engineering, establishing the viability of sequence-to-function deep learning models years before the present generation of transformer-based architectures [38]. Subsequent models, including Basenji and Enformer, extended this paradigm by incorporating dilated convolutions and, in Enformer's case, transformer self-attention layers to capture regulatory interactions across distances exceeding 100 kilobases, substantially improving gene expression prediction accuracy from sequence alone relative to purely convolutional predecessors [39].

7.2 Transformer-Based Genomic Models

The adaptation of transformer architectures, originally developed for natural language processing, to genomic sequences has produced a new class of genomic foundation models. DNABERT applied a BERT-style masked-language-modeling objective to k-mer-tokenized human genomic sequence, demonstrating that a single pretrained model could be fine-tuned to state-of-the-art performance on promoter prediction, splice-site prediction, and transcription factor binding site prediction with comparatively small amounts of task-specific labeled data [40]. DNABERT-2 subsequently replaced fixed k-mer tokenization with byte-pair encoding and expanded pretraining to multi-species genomic corpora, improving both computational efficiency and cross-species generalization, and introduced the Genome Understanding Evaluation benchmark to standardize comparison across genomic language models [41].

The Nucleotide Transformer family scaled this approach further, pretraining models ranging from 50 million to 2.5 billion parameters on genomes spanning thousands of human individuals and hundreds of additional species, and demonstrated that the resulting sequence representations alone, without fine-tuning, matched or outperformed specialized methods on a substantial majority of evaluated molecular phenotype prediction tasks [42]. More recent architectures, including HyenaDNA and the Evo genomic foundation model, have addressed the quadratic computational cost of self-attention by adopting long-convolution and state-space-inspired architectures capable of modeling single-nucleotide-resolution context windows extending to a million base pairs or more, substantially exceeding the practical context length of standard transformer architectures and enabling sequence design and generation tasks beyond simple classification [43], [44].

7.3 AI-Assisted Annotation and Pipeline Integration

Beyond standalone prediction tasks, machine learning is increasingly embedded within conventional bioinformatics pipelines rather than replacing them outright. Functional annotation pipelines now routinely incorporate learned models to prioritize candidate regulatory variants, predict the pathogenicity of missense mutations, and impute the functional consequence of non-coding variation that traditional annotation databases, dependent on direct experimental characterization, cannot otherwise classify [45]. The broader methodological shift documented across this literature is one from models that learn statistical relationships among sequencing reads, the domain of classical aligners and variant callers, toward models that learn the regulatory grammar of the genome itself, a transition with substantial implications for how future sequence analysis pipelines will be architected [45], [46].

It is important to note, however, that genomic language models remain an active and rapidly evolving research area rather than a fully mature replacement for established pipelines. Independent benchmarking has reported mixed performance for genomic foundation models on certain core tasks relative to specialized classical tools, particularly tasks requiring precise base-level alignment or reasoning that depends on integration with external reference information rather than sequence pattern recognition alone, indicating that the most plausible near-term trajectory is hybrid pipelines in which learned models complement, rather than supplant, classical alignment- and statistics-based methods [47].

8. Challenges and Limitations

8.1 Big Genomic Data

The scale of genomic data generation has long outpaced even the historically rapid growth rates associated with other big-data domains; one influential analysis projected that genomic data acquisition could, within a decade, rival or exceed the data demands of astronomy and other classically data-intensive sciences, driven by the falling cost of sequencing relative to the cost of data storage and transfer [16]. Population-scale sequencing initiatives now routinely generate multi-petabyte datasets, and the computational infrastructure required to store, transfer, and process such volumes increasingly exceeds the practical capacity of individual academic computing clusters, motivating the broader migration toward cloud computing platforms discussed in Section 9.

8.2 Computational Resource Requirements

Memory and processor requirements vary substantially across the tools reviewed in this paper, from the modest few-gigabyte footprint of Minimap2 or Salmon to the considerably larger requirements of STAR's suffix-array index, and this variation has direct practical consequences for tool accessibility, particularly in resource-limited research settings such as undergraduate laboratories or institutions in low- and middle-income countries. The increasing prevalence of deep learning-based tools compounds this constraint, as training and, to a lesser extent, inference for transformer-based genomic models typically require graphics processing unit acceleration, a resource not uniformly available across the global research community.

8.3 Data Quality Issues

Sequencing data quality varies substantially across platforms, library preparation protocols, and even individual sequencing runs, and systematic quality issues, including adapter contamination, GC-content bias, and platform-specific error signatures, propagate through downstream analysis if not adequately addressed during preprocessing [48]. Variant-calling benchmarking studies have repeatedly demonstrated that the choice of upstream quality control and alignment parameters can materially affect downstream variant call concordance, particularly for difficult genomic regions such as low-complexity repeats and segmental duplications [49].

8.4 Reproducibility Challenges

Reproducibility remains a recognized challenge in computational genomics, attributable to the combinatorial space of tool versions, parameter settings, reference genome builds, and annotation versions used across studies, a problem partially addressed by the adoption of workflow management systems and containerization, which encapsulate complete pipeline specifications for redistribution and exact re-execution [50]. Best-practice frameworks for variant calling in particular now explicitly recommend standardized, versioned pipelines precisely because subtle parameter differences between nominally equivalent pipelines have been shown to produce materially different variant call sets in clinical contexts [49].

8.5 Ethical Considerations

The increasing clinical and consumer availability of genomic sequencing raises ethical considerations that extend beyond the purely computational scope of this review but bear directly on responsible tool development and deployment. Genomic data, unlike most other biomedical data types, is inherently identifying and carries information relevant to biological relatives who have not themselves consented to data collection, raising distinctive privacy and consent challenges that computational pipelines and data-sharing infrastructures must address through technical safeguards such as access-controlled repositories and differential privacy mechanisms, alongside the policy frameworks that govern their use [9].

9. Future Directions

9.1 Cloud Bioinformatics

Cloud computing platforms have progressively become the default execution environment for large-scale genomic analysis, offering elastic compute and storage capacity that scales with project demand rather than requiring fixed institutional infrastructure investment [51]. Platforms such as Terra, DNAnexus, and the cloud-native implementations of GATK best-practice pipelines exemplify this shift, and the broader trend toward containerized, workflow-language-defined pipelines, expressed in standards such as the Workflow Description Language or Common Workflow Language, reflects an industry-wide convergence on portable, cloud-deployable analysis specifications that simultaneously address scalability and reproducibility concerns [51], [52].

9.2 AI-Driven Sequence Analysis

The trajectory of genomic foundation models suggests that future sequence analysis pipelines will increasingly incorporate pretrained, fine-tunable models as standard components alongside classical alignment and statistical tools, particularly for tasks involving regulatory element prediction, variant effect prediction, and the interpretation of non-coding genomic regions that have historically resisted rule-based annotation [42], [44]. The continuing reduction in computational cost of transformer inference, alongside the development of long-context architectures capable of modeling regulatory interactions across substantial genomic distances, positions these models to address biological questions, such as the cumulative regulatory effect of distal variants, that were previously intractable for purely alignment-based approaches [39].

9.3 Multi-Omics Integration

A clear trend across recent genomics literature is the move from single-modality analysis, DNA sequence alone or RNA expression alone, toward integrated multi-omics frameworks that jointly model genomic, transcriptomic, epigenomic, and proteomic data within a unified analytical structure. Single-cell multi-omics technologies, capable of measuring chromatin accessibility, gene expression, and, in some platforms, protein abundance within the same individual cell, have driven substantial methodological development in tools such as Seurat and Scanpy, which now support joint embedding and integration of multiple data modalities rather than treating each as an independent analysis [52], [53].

9.4 Precision Medicine

The convergence of faster, more accurate sequencing technologies, deep learning-augmented variant interpretation, and cloud-scale computational infrastructure is widely regarded as a principal enabler of precision medicine, in which clinical decisions are informed directly by an individual patient's genomic profile [9]. Realizing this potential at scale will require continued improvement in the speed and reliability of clinical-grade variant calling pipelines, alongside the development of standardized, interpretable reporting frameworks capable of translating raw variant calls into clinically actionable findings within timeframes compatible with clinical decision-making.

9.5 Real-Time Genomic Analytics

Nanopore sequencing's capacity for real-time data streaming, in which basecalling and downstream analysis can proceed concurrently with sequencing rather than only after a run completes, has opened a distinct frontier of real-time genomic analytics, with demonstrated applications in rapid pathogen identification during outbreak response and adaptive sampling strategies that dynamically enrich for regions of interest during sequencing itself [14]. Continued development of streaming-compatible alignment and variant-calling algorithms, capable of processing data incrementally rather than requiring complete datasets, is likely to extend this real-time paradigm to a broader range of diagnostic and surveillance applications in the coming years.

10. Conclusion

This review has examined the algorithmic foundations, computational trade-offs, and practical applications of the principal DNA and RNA sequence analysis tools in current bioinformatics practice. Several conclusions emerge from this comparative synthesis. First, no single tool dominates across all use cases; rather, tool selection should be governed by data modality, read length and error profile, available computational resources, and the specific biological question under investigation. Researchers working with short, accurate Illumina genomic reads are well served by BWA-MEM or Bowtie2, while those working with long, error-prone third-generation reads should default to Minimap2, which has effectively become the standard aligner for that data type. RNA-seq researchers face a more consequential methodological choice between alignment-based pipelines, exemplified by STAR paired with StringTie, which offer superior sensitivity for novel isoform and junction discovery at substantial memory cost, and alignment-free quantification via Salmon or Kallisto, which offer dramatically faster runtime and lower resource requirements at the cost of dependency on transcriptome reference completeness.

Second, the historical trajectory of this field, from exhaustive dynamic programming to heuristic index-based alignment to, most recently, learned sequence representations, reflects a consistent pattern in which algorithmic innovation has been driven primarily by the scaling demands of sequencing technology rather than by purely theoretical advances in computer science. This pattern suggests that the continuing evolution of sequencing platforms toward longer reads, higher throughput, and real-time data delivery will continue to reshape the computational requirements, and therefore the dominant algorithmic paradigms, of sequence analysis software for the foreseeable future.

Third, while deep learning and transformer-based genomic models have demonstrated substantial promise, particularly for variant calling and regulatory sequence interpretation, the evidence reviewed here does not support treating these models as wholesale replacements for established classical tools. The most defensible near-term strategy for practicing researchers, and the one implicitly adopted across the recent literature, is a hybrid approach in which classical alignment and statistical tools remain the backbone of standard pipelines, supplemented selectively by learned models for tasks, such as non-coding variant interpretation, for which classical methods offer limited resolution.

For students and early-career researchers selecting tools for a new project, three practical recommendations follow directly from this review. Tool choice should be validated against the specific data type and downstream analytical goal of the study, rather than selected by default or convention alone, since the comparative tables presented in Section 6 demonstrate that performance characteristics vary meaningfully across tools even within a single category. Computational resource constraints should be assessed before committing to memory-intensive tools such as STAR, since lower-memory alternatives such as HISAT2 or Salmon can deliver comparable or, for certain tasks, superior practical utility at substantially reduced infrastructure cost. Finally, researchers should track the rapidly evolving genomic foundation model literature as a complementary rather than substitutive resource, given the strong likelihood that hybrid classical-deep learning pipelines will constitute standard bioinformatics practice within the coming several years.

References

[1] S. F. Altschul, W. Gish, W. Miller, E. W. Myers, and D. J. Lipman, "Basic local alignment search tool," J. Mol. Biol., vol. 215, no. 3, pp. 403-410, 1990.

[2] C. Camacho et al., "BLAST+: architecture and applications," BMC Bioinformatics, vol. 10, p. 421, 2009.

[3] H. Li and R. Durbin, "Fast and accurate short read alignment with Burrows-Wheeler transform," Bioinformatics, vol. 25, no. 14, pp. 1754-1760, 2009.

[4] H. Li, "Aligning sequence reads, clone sequences and assembly contigs with BWA-MEM," arXiv:1303.3997, 2013.

[5] B. Langmead and S. L. Salzberg, "Fast gapped-read alignment with Bowtie 2," Nat. Methods, vol. 9, no. 4, pp. 357-359, 2012.

[6] D. Kim, J. M. Paggi, C. Park, C. Bennett, and S. L. Salzberg, "Graph-based genome alignment and genotyping with HISAT2 and HISAT-genotype," Nat. Biotechnol., vol. 37, no. 8, pp. 907-915, 2019.

[7] H. Li, "Minimap2: pairwise alignment for nucleotide sequences," Bioinformatics, vol. 34, no. 18, pp. 3094-3100, 2018.

[8] H. Li, "New strategies to improve minimap2 alignment accuracy," Bioinformatics, vol. 37, no. 23, pp. 4572-4574, 2021.

[9] A. Dobin et al., "STAR: ultrafast universal RNA-seq aligner," Bioinformatics, vol. 29, no. 1, pp. 15-21, 2013.

[10] R. Patro, G. Duggal, M. I. Love, R. A. Irizarry, and C. Kingsford, "Salmon provides fast and bias-aware quantification of transcript expression," Nat. Methods, vol. 14, no. 4, pp. 417-419, 2017.

[11] N. L. Bray, H. Pimentel, P. Melsted, and L. Pachter, "Near-optimal probabilistic RNA-seq quantification with kallisto," Nat. Biotechnol., vol. 34, no. 5, pp. 525-527, 2016.

[12] M. Pertea et al., "StringTie enables improved reconstruction of a transcriptome from RNA-seq reads," Nat. Biotechnol., vol. 33, no. 3, pp. 290-295, 2015.

[13] S. Kovaka et al., "Transcriptome assembly from long-read RNA-seq alignments with StringTie2," Genome Biol., vol. 20, p. 278, 2019.

[14] C. Trapnell et al., "Transcript assembly and quantification by RNA-Seq reveals unannotated transcripts and isoform switching during cell differentiation," Nat. Biotechnol., vol. 28, no. 5, pp. 511-515, 2010.

[15] C. Trapnell et al., "Differential analysis of gene regulation at transcript resolution with RNA-seq," Nat. Biotechnol., vol. 31, no. 1, pp. 46-53, 2013.

[16] R. Poplin et al., "A universal SNP and small-indel variant caller using deep neural networks," Nat. Biotechnol., vol. 36, no. 10, pp. 983-987, 2018.

[17] A. McKenna et al., "The Genome Analysis Toolkit: a MapReduce framework for analyzing next-generation DNA sequencing data," Genome Res., vol. 20, no. 9, pp. 1297-1303, 2010.

[18] T. Yun, H. Li, P.-C. Chang, M. F. Lin, A. Carroll, and C. Y. McLean, "Accurate, scalable cohort variant calls using DeepVariant and GLnexus," Bioinformatics, vol. 36, no. 24, pp. 5582-5589, 2021.

[19] H. Li et al., "The Sequence Alignment/Map format and SAMtools," Bioinformatics, vol. 25, no. 16, pp. 2078-2079, 2009.

[20] P. Danecek et al., "Twelve years of SAMtools and BCFtools," Gigascience, vol. 10, no. 2, p. giab008, 2021.

[21] Y. Ji, Z. Zhou, H. Liu, and R. V. Davuluri, "DNABERT: pre-trained Bidirectional Encoder Representations from Transformers model for DNA-language in genome," Bioinformatics, vol. 37, no. 15, pp. 2112-2120, 2021.

[22] Z. Zhou et al., "DNABERT-2: efficient foundation model and benchmark for multi-species genomes," in Proc. Int. Conf. Learning Representations (ICLR), 2024.

[23] H. Dalla-Torre et al., "Nucleotide Transformer: building and evaluating robust foundation models for human genomics," Nat. Methods, vol. 22, pp. 287-297, 2025.

[24] Z. Avsec et al., "Effective gene expression prediction from sequence by integrating long-range interactions," Nat. Methods, vol. 18, no. 10, pp. 1196-1203, 2021.

[25] E. Nguyen et al., "HyenaDNA: long-range genomic sequence modeling at single nucleotide resolution," in Proc. NeurIPS, 2023.

[26] E. Nguyen et al., "Sequence modeling and design from molecular to genome scale with Evo," Science, vol. 386, p. eado9336, 2024.

[27] J. Jumper et al., "Highly accurate protein structure prediction with AlphaFold," Nature, vol. 596, no. 7873, pp. 583-589, 2021.

[28] A. M. Wenger et al., "Accurate circular consensus long-read sequencing improves variant detection and assembly of a human genome," Nat. Biotechnol., vol. 37, no. 10, pp. 1155-1162, 2019.

[29] R. R. Wick, L. M. Judd, C. L. Gorrie, and K. E. Holt, "Unicycler: resolving bacterial genome assemblies from short and long sequencing reads," PLoS Comput. Biol., vol. 13, no. 6, p. e1005595, 2017.

[30] M. Kolmogorov, J. Yuan, Y. Lin, and P. A. Pevzner, "Assembly of long, error-prone reads using repeat graphs," Nat. Biotechnol., vol. 37, no. 5, pp. 540-546, 2019.

[31] H. Cheng, G. T. Concepcion, X. Feng, H. Zhang, and H. Li, "Haplotype-resolved de novo assembly using phased assembly graphs with hifiasm," Nat. Methods, vol. 18, no. 2, pp. 170-175, 2021.

[32] A. Bankevich et al., "SPAdes: a new genome assembly algorithm and its applications to single-cell sequencing," J. Comput. Biol., vol. 19, no. 5, pp. 455-477, 2012.

[33] M. G. Grabherr et al., "Full-length transcriptome assembly from RNA-Seq data without a reference genome," Nat. Biotechnol., vol. 29, no. 7, pp. 644-652, 2011.

[34] A. Frankish et al., "GENCODE 2021," Nucleic Acids Res., vol. 49, no. D1, pp. D916-D923, 2021.

[35] T. Stuart et al., "Comprehensive integration of single-cell data," Cell, vol. 177, no. 7, pp. 1888-1902, 2019.

[36] F. A. Wolf, P. Angerer, and F. J. Theis, "SCANPY: large-scale single-cell gene expression data analysis," Genome Biol., vol. 19, p. 15, 2018.

[37] C. Soneson, A. Srivastava, R. Patro, and M. B. Stadler, "Preprocessing choices affect RNA velocity results for droplet scRNA-seq data," PLoS Comput. Biol., vol. 17, no. 1, p. e1008585, 2021.

[38] V. Marx, "Method of the Year: long-read sequencing," Nat. Methods, vol. 20, no. 1, pp. 6-11, 2023.

[39] G. A. Logsdon, M. R. Vollger, and E. E. Eichler, "Long-read human genome sequencing and its applications," Nat. Rev. Genet., vol. 21, no. 10, pp. 597-614, 2020.

[40] S. L. Amarasinghe et al., "Opportunities and challenges in long-read sequencing data analysis," Genome Biol., vol. 21, p. 30, 2020.

[41] A. Conesa et al., "A survey of best practices for RNA-seq data analysis," Genome Biol., vol. 17, p. 13, 2016.

[42] Z. D. Stephens et al., "Big data: astronomical or genomical?," PLoS Biol., vol. 13, no. 7, p. e1002195, 2015.

[43] M. C. Schatz, B. Langmead, and S. L. Salzberg, "Cloud computing and the DNA data race," Nat. Biotechnol., vol. 28, no. 7, pp. 691-693, 2010.

[44] B. Langmead and A. Nellore, "Cloud computing for genomic data analysis and collaboration," Nat. Rev. Genet., vol. 19, no. 4, pp. 208-219, 2018.

[45] S. Goodwin, J. D. McPherson, and W. R. McCombie, "Coming of age: ten years of next-generation sequencing technologies," Nat. Rev. Genet., vol. 17, no. 6, pp. 333-351, 2016.

[46] S. Sandmann et al., "Evaluating variant calling tools for non-matched next-generation sequencing data," Sci. Rep., vol. 7, p. 43169, 2017.

[47] D. C. Koboldt, "Best practices for variant calling in clinical sequencing," Genome Med., vol. 12, p. 91, 2020.

[48] E. J. Topol, "High-performance medicine: the convergence of human and artificial intelligence," Nat. Med., vol. 25, no. 1, pp. 44-56, 2019.

[49] J. Zhou and O. G. Troyanskaya, "Predicting effects of noncoding variants with deep learning-based sequence model," Nat. Methods, vol. 12, no. 10, pp. 931-934, 2015.

[50] G. Eraslan, Z. Avsec, J. Gagneur, and F. J. Theis, "Deep learning: new computational modelling techniques for genomics," Nat. Rev. Genet., vol. 20, no. 7, pp. 389-403, 2019.

[51] S. R. Choi and M. Lee, "Transformer architecture and attention mechanisms in genome data analysis: a comprehensive review," Biology (Basel), vol. 12, no. 7, p. 1033, 2023.

[52] T. Hu, N. Chitnis, D. Monos, and A. Dinh, "Next-generation sequencing technologies: an overview," Hum. Immunol., vol. 82, no. 11, pp. 801-811, 2021.

[53] J. M. Heather and B. Chain, "The sequence of sequencers: the history of sequencing DNA," Genomics, vol. 107, no. 1, pp. 1-8, 2016.

[54] Z. Wang, M. Gerstein, and M. Snyder, "RNA-Seq: a revolutionary tool for transcriptomics," Nat. Rev. Genet., vol. 10, no. 1, pp. 57-63, 2009.

[55] G. A. Van der Auwera and B. D. O'Connor, Genomics in the Cloud: Using Docker, GATK, and WDL in Terra. Sebastopol, CA, USA: O'Reilly Media, 2020.